

Document Question Answering Assistant (RAG) MVP

Retrieval-Augmented Generation (RAG) combines document retrieval with a large language model to generate accurate answers. It reduces hallucinations and uses external knowledge without fine-tuning.

FAISS (Facebook AI Similarity Search) is a library for efficient similarity search and clustering of dense vectors.

Sentence Transformers is a library that generates state-of-the-art embeddings for sentences, paragraphs, and images. The all-MiniLM-L6-v2 model is a popular choice for fast and accurate embeddings.

Groq provides ultra-fast inference for large language models using its Tensor Streaming Processor (TSP) or Language Processing Unit (LPU) architecture, enabling real-time QA systems.